

Shaun Ting

(567) 901-1716 | shaunting.ck@gmail.com | [linkedin.com/in/shaunting](https://www.linkedin.com/in/shaunting) | github.com/shaunting

EDUCATION

University of Minnesota

Master of Science in Data Science

GPA: 4.0

Aug. 2025 – May 2027

University of Toledo

Bachelor of Science in Data Science, Minor in Computer Science

GPA: 3.7

Jan. 2022 – Dec. 2023

EXPERIENCE

AI Engineer

uNote

Aug. 2025 – Present

Minneapolis, MN

- Built a **Retrieval-Augmented Generation (RAG)** pipeline in **Python (FastAPI)** that extracts course PDFs into **embeddings**, storing them in a **vector database (pgvector)** for semantic search under 500ms.
- Designed **database schemas** for 6 core modules and built 10+ **REST APIs** in **Python (Django)** for courses, assignments, chatbot, interaction, files, and user management.
- Engineered a multi-**LLM** routing layer with **streaming responses** integrating **OpenAI**, **Anthropic**, and **Gemini APIs** with configurable settings, conversation memory, and answer controls for a personalized AI tutor.
- Deployed the platform using **Docker** and **AWS S3** with **CI/CD** through **GitHub Actions**, supporting 16,000+ concurrent users.

Senior Data Analyst

United Way of Rhode Island

Feb. 2024 – Feb. 2026

Providence, RI

- Built scheduled **ETL pipelines** in **Python** and **SQL** to clean, validate, and transfer donor data across 6 departments, reducing manual reporting time and standardizing cross-departmental data access.
- Integrated the **Mailchimp API** using **Python (Requests)** to pull campaign performance data, generating insights on donor engagement that informed marketing strategy.
- Analyzed a 100K-row donor dataset using **Python (pandas, NumPy, Matplotlib)** to identify retention patterns and acquisition sources, presenting findings to C-suite to inform \$13M in philanthropic decisions.
- Designed **ETL pipelines** in **SQL** and **PySpark** with **dimensional modeling** on **Microsoft Fabric**, building the data warehouse layer that powered real-time **Power BI** dashboards across the organization.

Data Science Intern

Department of Children, Youth & Families

June 2023 – Sep. 2023

Providence, RI

- Built **classification models (logistic regression)** in **SAS** to quantify racial disparities in child welfare outcomes, presenting feature-level insights to Child Protective Services supervisors to inform policy review.
- Automated **ETL pipelines** in **SAS** to clean, de-duplicate, and merge investigation and removal datasets across 9,000+ child welfare cases, reducing data preparation time by 18%.

PROJECTS

LLM Adversarial Robustness in Educational AI

- Designed adversarial **prompt engineering** attacks against two open-weight **LLMs** using **Python**, **LangChain**, and **HuggingFace**, evaluating system prompt compliance across 180 outputs using the **OpenAI API** as judge.

AI Clinical Documentation System

- Built a **Flask ETL pipeline** that transcribes doctor-patient audio via **Whisper API**, extracts structured notes using **Gemini API**, and stores results in a normalized **PostgreSQL** schema.

Semi-Supervised Floor Plan Semantic Segmentation

- Built a **semi-supervised computer vision** pipeline in **Python (PyTorch)** using FixMatch, training **CNN** architectures to achieve near fully-supervised segmentation accuracy with only 15% labeled data.

Fitness Workout Tracker

- Built a **classification** pipeline in **Python (Scikit-learn, Pandas, NumPy)** on sensor data, engineering 100+ features with **PCA**, **Fourier transforms**, and **K-Means** to achieve 99% accuracy with **Random Forest**.

TECHNICAL SKILLS

Languages: Python, SQL, R, SAS

AI / ML: LangChain, HuggingFace, OpenAI API, Gemini API, RAG, Embeddings, NLP, Prompt Engineering, Computer Vision, PyTorch, Scikit-learn

Backend: Django, FastAPI, Flask, REST APIs, PostgreSQL, ETL Pipelines, PySpark

Cloud / DevOps: AWS (S3), Microsoft Fabric, Docker, Git/GitHub, CI/CD (GitHub Actions)

Other: Agile/Scrum, Pandas, NumPy, Matplotlib, Power BI